

# SCTP as a First-Class Transport in Go on Linux: An Implementation and Interoperability Study

Olivier Van Acker  
Independent Researcher  
London, United Kingdom  
olivier@robotmotel.com

**Abstract**—This paper revisits a 2013 SCTP-in-Go implementation and re-establishes the design on Linux in a modern Go runtime tree. The central claim is that SCTP should be treated as a first-class transport protocol next to TCP and UDP, not as an exotic add-on. The argument is architectural and empirical: Internet protocol layering was designed to evolve, SCTP is a mature IETF transport, and practical integration in a contemporary runtime is straightforward when the networking stack is already protocol-agnostic at key extension points. We describe the Linux integration in Go’s `net` package, show API parity with existing TCP/UDP patterns, and evaluate interoperability against an independent C++ endpoint built on Linux SCTP APIs. Results show that the implementation passes in-tree SCTP tests and bilateral Go  $\leftrightarrow$  C++ interop scenarios while preserving SCTP-specific metadata such as stream identifier and PPID. The outcome supports a broader position: modern language runtimes should expose SCTP as a standard transport option for message-oriented and resilience-aware applications rather than forcing all designs through TCP or UDP.

**Index Terms**—SCTP, Go runtime, Linux networking, transport protocols, interoperability, sockets API

## I. INTRODUCTION

Transport protocol choice in many software stacks has narrowed to TCP or UDP, even though Internet architecture never required such a binary choice. The original internetworking design emphasized extensibility and evolution at protocol boundaries [1]–[3]. SCTP was standardized to provide reliable, congestion-controlled, message-oriented transport with features that are difficult or expensive to emulate correctly above TCP or UDP [4]–[6].

This paper revisits an earlier 2013 implementation effort and updates the work for Linux and a modern Go runtime source tree. The goal is not to present SCTP as “new”, but to demonstrate that it can be integrated into a mainstream language runtime with the same first-class treatment given to TCP and UDP. In this context, first-class means:

- explicit address and connection types in the standard networking package,
- normal resolver, dial, and listen flows,
- direct support for protocol-specific metadata and controls,
- test coverage and interoperability validation against an independent implementation.

The implementation in this repository adds SCTP support directly inside Go’s `net` package on Linux, including public API surface (`SCTPAddr`, `SCTPConn`,

`DialSCTP`, `ListenSCTP`), Linux data-path support for `sendmsg/recvmsg`, and targeted test and interop harnesses.

The paper’s central thesis is that SCTP belongs next to TCP and UDP in modern runtime APIs because:

- 1) Internet protocol architecture is intentionally extensible [2], [3].
- 2) SCTP is a mature standards-track transport with broad, production-relevant semantics [4], [5], [7], [8].
- 3) Runtime integration is technically modest when implemented through existing socket abstraction seams.
- 4) Real interoperability with C++ Linux SCTP endpoints demonstrates practical viability.

The remainder of this paper builds this argument from standards context, implementation design, and measured evaluation.

## II. WHY SCTP SHOULD BE FIRST-CLASS

### A. Architectural Basis: IP Was Designed to Evolve

The Internet stack has long separated network and transport concerns to permit independent evolution [1], [2], [9]. RFC 1958 makes this explicit: successful protocol design assumes change and adaptation over time [3]. Treating TCP and UDP as immutable endpoints of transport evolution is therefore an operational habit, not an architectural law.

### B. SCTP Is Not Experimental

SCTP is a standards-track transport protocol with a mature socket API and long-running implementations [4], [5], [10]. Relative to canonical TCP and UDP definitions [11], [12], it combines properties that many applications need:

- reliable, congestion-aware transport like TCP,
- preservation of message boundaries unlike TCP,
- multistreaming to reduce head-of-line coupling,
- optional multihoming and failover support at transport layer,
- protocol-level extensibility (e.g., partial reliability, schedulers).

These are not theoretical features; they are codified in the base protocol and extensions [7], [8], [13], [14].

### C. Industry Relevance

SCTP is also used as the transport substrate for WebRTC data channels [15], [16]. This matters for two reasons: first, it

disproves the claim that SCTP has no mainstream deployment; second, it shows that modern systems already depend on SCTP semantics when message orientation and stream multiplexing are required.

#### D. Comparison with TCP and UDP

Table I summarizes why SCTP deserves parity in language runtime APIs.

TABLE I  
TRANSPORT FEATURE COMPARISON

Feature	TCP	UDP	SCTP
Reliable delivery	Yes	No	Yes
Message boundaries	No	Yes	Yes
Built-in multistreaming	No	No	Yes
Transport multihoming	No	No	Yes
Congestion control	Yes	No	Yes
Socket API standardization	Yes	Yes	Yes

The practical conclusion is straightforward: if TCP and UDP are exposed natively in a runtime’s core networking API, SCTP can be justified by the same criterion used for those protocols, namely broad utility plus standards-based interoperability.

### III. DESIGN AND INTEGRATION IN THE GO RUNTIME

#### A. Integration Goal

The implementation target is parity with existing net package transport patterns: address resolution, dial/listen entry points, and connection methods should look and behave like TCP/UDP equivalents while exposing SCTP-specific controls where required.

#### B. Public API Additions

The Linux implementation introduces SCTP-specific types and functions in `src/net/sctpsock.go`, including:

- type SCTPAddr { IP net.IP; Port int; Zone string }
- ResolveSCTPAddr(network, address)
- DialSCTP(network, laddr, raddr)
- ListenSCTP(network, laddr)
- type SCTPConn with ReadFromSCTP, WriteToSCTP, SetNoDelay, SetInitOptions, SubscribeEvents

These additions mirror established Go networking style and maintain compatibility with common Conn/PacketConn usage.

#### C. Runtime Touchpoints

Implementation is localized to existing extensibility seams in net:

- Resolver and network parsing integration in `src/net/ipsock.go`
- Dial/listen dispatch integration in `src/net/dial.go`
- Linux SCTP socket operations in `src/net/sctpsock_posix.go`

- Linux SCTP constants and control-message marshaling in `src/net/sctpsock_linux.go`
- Stubbed behavior for unsupported targets in `src/net/sctpsock_stub.go`

This is the core evidence for integration ease: SCTP support was added by extending existing abstractions rather than redesigning the runtime networking architecture.

#### D. Data Path and Metadata

The implementation uses `SOCK_SEQPACKET` with `IPPROTO_SCTP`. Message transmission and reception are built on `sendmsg/recvmsg`, with ancillary control messages carrying SCTP metadata (`SCTP SNDINFO`, `SCTP RCVINFO`). This preserves message orientation while exposing stream and PPID metadata to applications.

#### E. One-to-Many First

The current design emphasizes one-to-many SCTP semantics for Linux. DialSCTP uses an unconnected socket model with remote destination retained for default sends, and ListenSCTP follows passive receive behavior on `SOCK_SEQPACKET`. This keeps the API simple while supporting key SCTP behavior.

#### Integration map in Go net package

- 1) Parse network: `sctp`, `sctp4`, `sctp6`
- 2) Resolve endpoint: `SCTPAddr`
- 3) Dial/listen dispatch in shared socket path
- 4) Linux SCTP options and cmsg handlers
- 5) Public `SCTPConn` methods for metadata-aware I/O

Fig. 1. SCTP follows existing TCP/UDP extension seams in Go networking internals.

### IV. EVALUATION METHOD

#### A. Objectives

The evaluation verifies two claims:

- 1) SCTP is integrated as a first-class runtime transport, not as an out-of-tree shim.
- 2) The implementation interoperates with an independent C++ SCTP endpoint on Linux.

#### B. Environment

Measurements in this paper were captured on:

- OS: Ubuntu 25.10, Linux kernel 6.17.0-12-generic
- Architecture: linux/amd64
- Go tree: in-repo build, version `go1.27-devel_c9cbeb0`
- SCTP kernel module enabled (`sctp`)

### C. Go Package Tests

Targeted tests in `net` were executed:

- `TestSCTPLoopbackReadWrite`
- `TestSCTPUnsupportedOnBadNetwork`
- `TestResolveSCTPAddrUnknownNetwork`
- `TestParseNetworkSCTP`
- `TestSCTPAddrString`

These tests validate local SCTP I/O behavior, error semantics, and network/address parsing paths.

### D. Go $\leftrightarrow$ C++ Interoperability Matrix

Interop uses the repository harness at `misc/sctp-interop/harness/run_matrix.sh` and covers:

- 1) Go server receiving from C++ client.
- 2) C++ server receiving from Go client.

Both scenarios verify payload integrity and SCTP metadata transport (stream ID and PPID). The C++ side uses Linux SCTP userspace API calls and `recvmsg` ancillary parsing.

### E. Timing Measurement

The complete matrix runner was executed 20 times. Wall-clock duration per run was recorded from shell timestamps. This provides a lightweight integration performance indicator for repeated bring-up, cross-language exchange, and teardown under the same configuration.

## V. RESULTS

### A. Go In-Tree SCTP Tests

All targeted Go `net` SCTP tests passed in a single run. The test run reported successful socket creation and loopback communication on `AF_INET/SOCK_SEQPACKET/IPPROTO_SCTP`.

### B. Interop Matrix Outcome

Both interoperability directions passed:

- Go server received C++ payload with expected stream and PPID.
- C++ server received Go payload with expected stream and PPID.

Table II summarizes the observed data-plane correctness.

TABLE II  
OBSERVED Go  $\leftrightarrow$  C++ SCTP INTEROPERABILITY

Scenario	Observed output
Go server $\leftarrow$ C++ client	stream=3 ppid=101 payload=cpp-to-go
C++ server $\leftarrow$ Go client	stream=4 ppid=202 payload=go-to-cpp

### C. Repeated Matrix Runtime

For 20 repeated runs of the full interop matrix:

- mean runtime: 2.187 037 s
- minimum runtime: 2.177 443 s
- maximum runtime: 2.203 883 s
- standard deviation: 0.006 289 s

The low variation suggests stable harness behavior and reproducible cross-language SCTP exchange in the tested environment.

### 20-run matrix timing summary

Mean: 2.187 s

Min/Max: 2.177 s / 2.204 s

Std. dev.: 0.006 s

Data source:

[paper/repro/data/interop\\_matrix\\_runtime\\_2026-02-14.csv](#)

Fig. 2. Runtime stability of the Go  $\leftrightarrow$  C++ SCTP matrix on Linux.

### D. Interpretation

The results support the implementation claim:

- 1) SCTP behaves as an ordinary transport option inside Go's standard networking paths.
- 2) SCTP-specific metadata survives end-to-end across independent implementations.
- 3) Integration and validation can be automated with the same workflow style used for other runtime networking features.

## VI. RELATED WORK

The original SCTP motivation and protocol model are well documented in standards and reference texts [4], [6]. The socket API formalization in RFC 6458 is directly relevant to runtime and language binding design [5]. Extensions such as partial reliability and stream scheduling further differentiate SCTP from TCP-centric abstractions [7], [8].

Several prior studies evaluated SCTP capabilities. Iyengar et al. explored concurrent multipath transfer over SCTP multihoming [17]. Penoff et al. described a portable userspace SCTP stack and portability/performance trade-offs [18]. These works support the view that SCTP is both implementable and practically useful across deployment contexts.

More generally, transport evolution research has shown how deployment friction can ossify protocol diversity [19]. This paper contributes from a runtime-engineering perspective: language-level first-class APIs materially reduce that friction by making non-TCP transports available through familiar interfaces.

Finally, WebRTC data channels demonstrate broad production exposure of SCTP semantics in modern applications [15], [16]. This real-world adoption weakens arguments that SCTP is niche or obsolete.

## VII. DISCUSSION

### A. Why This Supports “First-Class” Status

The implementation and evaluation provide concrete evidence for first-class SCTP treatment in a modern runtime:

- **API symmetry:** SCTP follows the same naming and usage patterns as TCP/UDP in Go’s `net` package.
- **Internal symmetry:** SCTP plugs into existing resolver, dial, listen, and file-descriptor workflows.
- **Interoperability:** wire-level behavior is compatible with independent Linux C++ SCTP endpoints.
- **Operational symmetry:** tests and harness execution integrate into standard CI-style workflows.

This is precisely what “first-class” should mean in a runtime context.

### B. The Extensibility Argument

Internet architecture intentionally supports protocol evolution [1]–[3]. Treating SCTP as a default-capable option is consistent with this model and avoids unnecessary application-layer reimplementation of transport semantics. This is aligned with end-to-end design reasoning: semantics that belong at transport should not be rebuilt repeatedly in each application [20]. In practice, forcing message-oriented or multi-stream designs onto TCP often reintroduces complexity that SCTP already standardizes.

### C. Additional Practical Arguments

Beyond architectural correctness, several practical arguments support broader SCTP exposure:

- **Correctness:** protocol-level message framing reduces application parsing ambiguity.
- **Performance isolation:** multistreaming can reduce cross-flow head-of-line coupling.
- **Resilience:** multihoming and association semantics improve failover options.
- **Security posture:** SCTP’s association handshake and extension mechanisms provide robust baseline behavior [4].

None of these claims require replacing TCP/UDP; they require giving developers a native, standards-based choice.

## VIII. LIMITATIONS AND THREATS TO VALIDITY

This study has clear scope boundaries:

- Evaluation was performed on one Linux environment and loopback/local-host paths.
- Interop matrix validates bidirectional Go  $\leftrightarrow$  C++ exchange, but does not exhaust all SCTP extensions.
- The current harness does not benchmark high-throughput bulk transfer against tuned TCP/UDP baselines.
- Multihoming failover scenarios were not fully exercised in a multi-interface fault-injection lab.

These limitations do not invalidate the central result (integration feasibility plus interop correctness), but they constrain generalization of quantitative performance conclusions.

Future work should include multi-host and multi-path experiments, scheduler/interleaving extension coverage, and comparative workload benchmarks for message-centric applications.

## IX. CONCLUSION

This paper presented a Linux-focused SCTP integration into a modern Go runtime and evaluated it with in-tree tests plus Go  $\leftrightarrow$  C++ interoperability scenarios. The evidence supports the main claim: SCTP can and should be treated as a first-class transport next to TCP and UDP in runtime networking APIs.

The broader point is architectural. Internet protocol design expects extensibility, and SCTP is a mature transport standard that addresses real application needs without forcing ad hoc reimplementation at higher layers. The implementation in this repository demonstrates that enabling SCTP in a mainstream runtime is not a disruptive redesign; it is an incremental, testable extension of existing networking abstractions.

Treating SCTP as first-class is therefore both technically justified and operationally practical.

## REFERENCES

- [1] V. Cerf and R. Kahn, “A protocol for packet network intercommunication,” *IEEE Transactions on Communications*, vol. 22, no. 5, pp. 637–648, 1974.
- [2] R. Braden, “Requirements for internet hosts – communication layers,” RFC Editor, RFC 1122, Oct. 1989. [Online]. Available: <https://www.rfc-editor.org/rfc/rfc1122>
- [3] B. Carpenter, “Architectural principles of the internet,” RFC Editor, RFC 1958, Jun. 1996. [Online]. Available: <https://www.rfc-editor.org/rfc/rfc1958>
- [4] R. Stewart, “Stream control transmission protocol,” RFC Editor, RFC 9260, Jun. 2022. [Online]. Available: <https://www.rfc-editor.org/rfc/rfc9260>
- [5] R. Stewart, M. Tuexen, P. Lei, and L. Ong, “Sockets api extensions for stream control transmission protocol (sctp),” RFC Editor, RFC 6458, Dec. 2011. [Online]. Available: <https://www.rfc-editor.org/rfc/rfc6458>
- [6] R. R. Stewart and Q. Xie, *Stream Control Transmission Protocol (SCTP): A Reference Guide*. Addison-Wesley, 2001.
- [7] R. Stewart, M. Ramalho, Q. Xie, M. Tuexen, and P. Conrad, “Stream control transmission protocol (sctp) partial reliability extension,” RFC Editor, RFC 3758, May 2004. [Online]. Available: <https://www.rfc-editor.org/rfc/rfc3758>
- [8] R. Stewart, M. Tuexen, I. Ruengeler, and S. Loreto, “Stream schedulers and user message interleaving for the stream control transmission protocol,” RFC Editor, RFC 8260, Nov. 2017. [Online]. Available: <https://www.rfc-editor.org/rfc/rfc8260>
- [9] J. Postel, “Internet protocol,” RFC Editor, RFC 791, Sep. 1981. [Online]. Available: <https://www.rfc-editor.org/rfc/rfc791>
- [10] M. Kerrisk, “sctp(7) – linux manual page,” <https://man7.org/linux/man-pages/man7/sctp.7.html>, 2025, accessed: 2026-02-14.
- [11] W. Eddy, “Transmission control protocol (tcp),” RFC Editor, RFC 9293, Aug. 2022. [Online]. Available: <https://www.rfc-editor.org/rfc/rfc9293>
- [12] J. Postel, “User datagram protocol,” RFC Editor, RFC 768, Aug. 1980. [Online]. Available: <https://www.rfc-editor.org/rfc/rfc768>
- [13] R. Stewart, P. Lei, M. Tuexen, and E. Rescorla, “Stream control transmission protocol (sctp) dynamic address reconfiguration,” RFC Editor, RFC 5061, Sep. 2007. [Online]. Available: <https://www.rfc-editor.org/rfc/rfc5061>
- [14] M. Tuexen, I. Ruengeler, P. Natarajan, M. Stewart, and M. Kuehlewind, “Sctp ndata: Stream schedulers and user message interleaving,” RFC Editor, RFC 7053, Dec. 2013. [Online]. Available: <https://www.rfc-editor.org/rfc/rfc7053>
- [15] R. Jesup, S. Loreto, and M. Tuexen, “WebRTC data channels,” RFC Editor, RFC 8831, Jan. 2021. [Online]. Available: <https://www.rfc-editor.org/rfc/rfc8831>

- [16] ——, “WebRTC data channel establishment protocol,” RFC Editor, RFC 8832, Jan. 2021. [Online]. Available: <https://www.rfc-editor.org/rfc/rfc8832>
- [17] J. R. Iyengar, P. D. Amer, and R. Stewart, “Concurrent multipath transfer using sctp multihoming over independent end-to-end paths,” *IEEE/ACM Transactions on Networking*, vol. 14, no. 5, pp. 951–964, 2006.
- [18] B. Penoff, A. Wagner, M. Tuexen, and I. Ruengeler, “Portable and performant user space sctp stack,” in *Proceedings of the 21st International Conference on Computer Communications and Networks (ICCCN)*, 2012, pp. 1–9.
- [19] M. Honda, Y. Nishida, C. Raiciu, A. Greenhalgh, M. Handley, and H. Tokuda, “Is it still possible to extend tcp?” in *Proceedings of the 11th ACM SIGCOMM Conference on Internet Measurement (IMC)*, 2011, pp. 181–194.
- [20] J. H. Saltzer, D. P. Reed, and D. D. Clark, “End-to-end arguments in system design,” *ACM Transactions on Computer Systems*, vol. 2, no. 4, pp. 277–288, 1984.